

## Advanced Data Analytics using Hive of Titanic: Machine Learning from Disaster dataset

### Calculate how many events occurred according to the event types

First, ran the below HiveQL query to calculate how many events occurred according to the event types of each driver.

```
SELECT driverid, eventType, count(*) FROM csc534.tdend2_truck_event GROUP BY driverId, eventType;
```

```
tdend2@node00:~
hive> SELECT driverid, eventType, count(*) FROM csc534.tdend2_truck_event GROUP BY driverId, eventType;
Query ID = tdend2_20241009170501_963f7c3b-1fc0-4290-b77f-148b6b584c20
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-10-09 17:05:14,853 Stage-1 map = 0%, reduce = 0%
2024-10-09 17:05:22,064 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.74 sec
2024-10-09 17:05:29,244 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 8.14 sec
MapReduce Total cumulative CPU time: 8 seconds 140 msec
Ended Job = job_1722897143033_1257
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.14 sec HDFS Read: 2283649 HDFS Write: 1610 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 140 msec
OK
10      Normal  727
10      Overspeed      2
10      Unsafe tail distance  1
11      Lane Departure  4
11      Normal  676
11      Overspeed      1
```

Last half of output:

```
tdend2@node00:~  
11 Lane Departure 4  
11 Normal 676  
11 Overspeed 1  
11 Unsafe following distance 1  
11 Unsafe tail distance 4  
12 Normal 745  
12 Unsafe following distance 1  
13 Lane Departure 1  
13 Normal 745  
14 Normal 745  
14 Unsafe following distance 1  
15 Lane Departure 1  
15 Normal 740  
16 Lane Departure 1  
16 Normal 744  
17 Lane Departure 1  
17 Normal 747  
18 Normal 750  
18 Overspeed 1  
19 Normal 744  
19 Unsafe following distance 1  
20 Normal 745  
20 Overspeed 1  
21 Normal 747  
21 Unsafe tail distance 1  
22 Normal 748  
22 Unsafe tail distance 1  
23 Lane Departure 1  
23 Normal 744  
24 Lane Departure 1  
24 Normal 744  
25 Normal 742  
25 Overspeed 1  
26 Normal 743  
26 Overspeed 1  
27 Normal 741  
27 Unsafe following distance 1  
28 Normal 743  
28 Overspeed 1  
29 Normal 749  
29 Overspeed 1  
30 Normal 740  
30 Unsafe following distance 1  
31 Lane Departure 1  
31 Normal 746  
32 Normal 746  
32 Unsafe following distance 1  
Time taken: 28.856 seconds, Fetched: 50 row(s)  
hive>
```

To show header or column name in the outputs, ran below command:

```
set hive.cli.print.header=true;  
set hive.resultset.use.unique.column.names=false;
```

### Calculate the total events that occurred

Next, ran the below HiveQL query to calculate the total events that occurred per driver.

SELECT driverId, count(\*) AS total\_events FROM csc534.tdend2\_truck\_event GROUP BY driverId;

```
tdend2@node00:~
hive> set hive.cli.print.header=true;
hive> SELECT driverId, count(*) AS total_events FROM csc534.tdend2_truck_event GROUP BY driverId;
Query ID = tdend2_20241009171433_27af75c6-b3c8-4799-83c3-9454086d39c8
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
24/10/09 17:14:33 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
24/10/09 17:14:33 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
Starting Job = job_1722897143033_1258, Tracking URL = http://node00.sun:8088/proxy/application_1722897143033_1258/
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/lib/hadoop/bin/hadoop job -kill job_1722897143033_1258
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-10-09 17:14:43,198 Stage-1 map = 0%, reduce = 0%
2024-10-09 17:14:51,396 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.66 sec
2024-10-09 17:14:59,583 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.07 sec
MapReduce Total cumulative CPU time: 7 seconds 70 msec
Ended Job = job_1722897143033_1258
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.07 sec HDFS Read: 2283400 HDFS Write: 524 HD
Total MapReduce CPU Time Spent: 7 seconds 70 msec
OK
driverid      total_events
10            730
11            686
12            746
13            746
```

Total events based on driver id with header :

```
tdend2@node00:~
driverid      total_events
10            730
11            686
12            746
13            746
14            746
15            741
16            745
17            748
18            751
19            745
20            746
21            748
22            749
23            745
24            745
25            743
26            744
27            742
28            744
29            750
30            741
31            747
32            747
Time taken: 27.223 seconds, Fetched: 23 row(s)
hive>
```

### Connect both GROUP BY result sets

Finally, tried to connect both GROUP BY result sets with UNION ALL as below:

Also need to modify the queries to the schema of both sides of the union to be matched

```
SELECT driverid, eventType, count(*) AS total FROM csc534.tdend2_truck_event GROUP BY driverid, eventType UNION ALL SELECT driverid, null as eventType, count(*) AS total FROM csc534.tdend2_truck_event GROUP BY driverid;
```

```
tdend2@node00:~  
hive> SELECT driverid, eventType, count(*) AS total FROM csc534.tdend2_truck_event GROUP BY driverid, eventType UNION ALL SELECT driverid, null as eventType, count(*) AS total FROM csc534.tdend2_truck_event GROUP BY driverid;  
Query ID = tdend2_20241009172041_45694824-c8ea-40b0-82b8-af85071859b2  
Total jobs = 3  
Launching Job 1 out of 3  
Number of reduce tasks not specified. Estimated from input data size: 1
```

```
tdend2@node00:~  
Stage-Stage-2: Map: 2 Cumulative CPU: 5.89 sec HDFS Read: 10396 HDFS Write: 2203 HDFS EC Read: 0 SUCCESS  
Total MapReduce CPU Time Spent: 20 seconds 130 msec  
OK  
_u1.driverid _u1.eventtype _u1.total  
10 Normal 727  
10 Overspeed 2  
10 Unsafe tail distance 1  
11 Lane Departure 4  
11 Normal 676  
11 Overspeed 1  
11 Unsafe following distance 1  
11 Unsafe tail distance 4  
12 Normal 745  
12 Unsafe following distance 1  
12 Lane Departure 1
```

Last half of output:

```
tdend2@node00:~  
10      NULL      730  
11      NULL      686  
12      NULL      746  
13      NULL      746  
14      NULL      746  
15      NULL      741  
16      NULL      745  
17      NULL      748  
18      NULL      751  
19      NULL      745  
20      NULL      746  
21      NULL      748  
22      NULL      749  
23      NULL      745  
24      NULL      745  
25      NULL      743  
26      NULL      744  
27      NULL      742  
28      NULL      744  
29      NULL      750  
30      NULL      741  
31      NULL      747  
32      NULL      747  
Time taken: 68.609 seconds, Fetched: 73 row(s)  
hive> _
```

It shows the driver's id, event type, and the total number of events from both query result sets. It took quite a long time compared to first query second query, because it read the datasets twice. If we need to analyze Big data, it could be a problem. This is where GROUPING SETS come in.

• (10pts) Section 2.2

```
SELECT driverId, eventType, count(*) AS occurrence  
FROM csc534.tdend2_truck_event  
GROUP BY driverId, eventType  
GROUPING SETS ((driverId, eventType), driverId);
```

```
tdend2@node00:~  
  
hive> SELECT driverId, eventType, count(*) AS occurrence  
      > FROM csc534.tdend2_truck_event  
      > GROUP BY driverId, eventType  
      > GROUPING SETS ((driverId, eventType), driverId);  
Query ID = tdend2_20241009172747_e42e4dac-04ce-40c7-bea3-1bab65c6b9eb
```

```

tdend2@node00:~
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.15 sec HDFS Read: 2284335 HDFS Write: 2116 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 150 msec
OK
driverid      eventtype      occurrence
10      NULL      730
10      Normal      727
10      Overspeed      2
10      Unsafe tail distance      1
11      NULL      686
11      Lane Departure      4
11      Normal      676
11      Overspeed      1
11      Unsafe following distance      1
11      Unsafe tail distance      4
12      NULL      746
12      Normal      745
12      Unsafe following distance      1
13      NULL      746

```

Last half of output:

```

29      Overspeed      1
30      NULL      741
30      Normal      740
30      Unsafe following distance      1
31      NULL      747
31      Lane Departure      1
31      Normal      746
32      NULL      747
32      Normal      746
32      Unsafe following distance      1
Time taken: 25.931 seconds, Fetched: 73 row(s)
hive>

```

We calculated how many events occurred according to each driver's event types and the total events that occurred per driver. Where the event column is null, we have the total sum of events that occurred to a driver across all event types.

### Analyzing Driver Risk factor

Analyzed the risks associated with drivers in detail. We will use event and driver's mileage records to calculate the driver's risk factor.

### Non-normal events and totals

First, filtered normal records from the events records (geolocation table), then created an unusual\_event table. Ran the below query.

```

CREATE TABLE csc534.tdend2_unusual_events
STORED AS ORC
AS
SELECT driverId, count(*) AS occurrence
FROM csc534.tdend2_truck_event
WHERE eventType != 'Normal'
GROUP BY driverId;

```

```

tdend2@node00:~
hive> CREATE TABLE csc534.tdend2_unusual_events
> STORED AS ORC
> AS
> SELECT driverId, count(*) AS occurrence
> FROM csc534.tdend2_truck_event
> WHERE eventType != 'Normal'
> GROUP BY driverId;
Query ID = tdend2_20241009173432_9ba86569-b44b-4c64-94c0-84dd27afb830
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
24/10/09 17:34:32 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
24/10/09 17:34:32 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
Starting Job = job_1722897143033_1268, Tracking URL = http://node00.sun:8088/proxy/application_1722897143033_1268
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/lib/hadoop/bin/hadoop job -kill job_1722897143033_1268
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-10-09 17:34:40,574 Stage-1 map = 0%, reduce = 0%
2024-10-09 17:34:46,712 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 6.13 sec
2024-10-09 17:34:55,901 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.45 sec
MapReduce Total cumulative CPU time: 10 seconds 450 msec
Ended Job = job_1722897143033_1268
Moving data to directory hdfs://node00.sun:8020/user/hive/warehouse/csc534.db/tdend2_unusual_events
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 10.45 sec HDFS Read: 2284022 HDFS Write: 365
Total MapReduce CPU Time Spent: 10 seconds 450 msec
OK
driverid      occurrence
Time taken: 25.813 seconds
hive>

```

The resulting table will count each driver's total unusual or abnormal events

```

SELECT *
FROM csc534.tdend2_unusual_events
LIMIT 5;

```

```

tdend2@node00:~
hive> SELECT *
> FROM csc534.tdend2_unusual_events
> LIMIT 5;
OK
tdend2_unusual_events.driverid tdend2_unusual_events.occurrence
10      3
11      10
12      1
13      1
14      1
Time taken: 0.077 seconds, Fetched: 5 row(s)
hive>

```

Also, created another table, totals, with total hours and miles.

```
CREATE TABLE csc534.tdend2_totals
STORED AS ORC
AS
SELECT driverId, sum(hours_logged) AS total_hours, sum(miles_logged) AS total_miles
FROM csc534.tdend2_timesheet
GROUP BY driverId;
```

```
tdend2@node00:~
hive> CREATE TABLE csc534.tdend2_totals
> STORED AS ORC
> AS
> SELECT driverId, sum(hours_logged) AS total_hours, sum(miles_logged) AS total_miles
> FROM csc534.tdend2_timesheet
> GROUP BY driverId;
Query ID = tdend2_20241009173842_da68db64-3a5f-4e5d-b927-b01d059bae42
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
24/10/09 17:38:42 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
24/10/09 17:38:42 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
Starting Job = job_1722897143033_1269, Tracking URL = http://node00.sun:8088/proxy/application_17
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/lib/hadoop/bin/hadoop job -
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-10-09 17:38:50,864 Stage-1 map = 0%, reduce = 0%
2024-10-09 17:38:59,049 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.72 sec
2024-10-09 17:39:05,178 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.14 sec
MapReduce Total cumulative CPU time: 6 seconds 140 msec
Ended Job = job_1722897143033_1269
Moving data to directory hdfs://node00.sun:8020/user/hive/warehouse/csc534.db/tdend2_totals
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.14 sec HDFS Read: 35166 HDFS Write: 630 HD
Total MapReduce CPU Time Spent: 6 seconds 140 msec
OK
driverid      total_hours    total_miles
Time taken: 24.635 seconds
hive>
```



We can check the resulting table as shown below:

```
SELECT *
FROM csc534.tdend2_totals
LIMIT 5;
```

```
tdend2@node00:~
hive> SELECT *
> FROM csc534.tdend2_totals
> LIMIT 5;
OK
tdend2_totals.driverid  tdend2_totals.total_hours  tdend2_totals.total_miles
10      3232      147150
11      3642      179300
12      2639      135962
13      2727      134126
14      2781      136624
Time taken: 0.066 seconds, Fetched: 5 row(s)
hive>
```

### 2.3.2 Perform JOIN operation

Performed a JOIN operation. The unusual\_events table has the driver's ID and count of their respective non-normal events. The totals table shows the hours and miles each driver travelled.

```
CREATE TABLE csc534.tdend2_joined
STORED AS ORC
AS
SELECT u.driverId, u.occurrence, t.total_hours, t.total_miles
FROM csc534.tdend2_unusual_events u
JOIN csc534.tdend2_totals t
ON (u.driverId = t.driverId);
```

```
tdend2@node00:~
hive> CREATE TABLE csc534.tdend2_joined
> STORED AS ORC
> AS
> SELECT u.driverId, u.occurrence, t.total_hours, t.total_miles
> FROM csc534.tdend2_unusual_events u
> JOIN csc534.tdend2_totals t
> ON (u.driverId = t.driverId);
Query ID = tdend2_20241009174402_5dba689a-bd52-4f00-875b-5c10b15de9dd
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/jars/log4j-core-2.17.0.jar:]
SLF4J: Found binding in [jar:file:/opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/jars/slf4j-api-1.7.25.jar:]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
```

```
tdend2@node00:~
Total MapReduce CPU Time Spent: 6 seconds 180 msec
OK
u.driverid      u.occurrence      t.total_hours      t.total_miles
Time taken: 30.668 seconds
hive>
```

The resulting data set will give us a driver's total miles and non-normal events.

```
SELECT *
FROM csc534.tdend2_joined
LIMIT 5;
```

```
tdend2@node00:~
hive> SELECT *
> FROM csc534.tdend2_joined
> LIMIT 5;
OK
tdend2_joined.driverid  tdend2_joined.occurrence  tdend2_joined.total_hours  tdend2_joined.total_miles
10      3      3232      147150
11      10     3642     179300
12      1     2639     135962
13      1     2727     134126
14      1     2781     136624
Time taken: 0.068 seconds, Fetched: 5 row(s)
hive>
```

### Compute driver risk factor

We will associate a driver risk factor with every driver. The risk factor for each driver is the number of unusual/abnormal occurrences over the total number of miles. A high number of unusual/abnormal occurrences over a short number of miles indicates high risk. Let's translate this intuition into an HiveQL query.

```
SELECT driverid, occurrence, total_hours, total_miles, total_miles/occurrence
AS riskfactor
FROM csc534.tdend2_joined
ORDER BY riskfactor ASC
LIMIT 5;
```

```
tdend2@node00:~
hive> SELECT driverid, occurrence, total_hours, total_miles, total_miles/occurrence
> AS riskfactor
> FROM csc534.tdend2_joined
> ORDER BY riskfactor ASC
> LIMIT 5;
Query ID = tdend2_20241009174824_f69294c6-8be2-49dd-8c03-04dd90577a95
```

```
tdend2@node00:~
Total MapReduce CPU Time Spent: 8 seconds 930 msec
OK
driverid  occurrence  total_hours  total_miles  riskfactor
11      10     3642     179300    17930.0
10      3      3232     147150    49050.0
13      1     2727     134126    134126.0
24      1     2647     134461    134461.0
20      1     2644     134564    134564.0
Time taken: 27.129 seconds, Fetched: 5 row(s)
hive>
```

The resulting data set will give us occurrence, total\_hours, total miles, total unusual/abnormal events, and what is a risk for a particular driver.

We exercised GROUPING SETS. There are additional ways to perform these aggregations, using CUBE or ROLLUP. These are almost like shortcuts. While CUBE returns all possible aggregation combinations, ROLLUP does it more hierarchically.

### Rollup

**Modify/rewrite the grouping set query using ROLLUP. Run and show the results. Explain the presentation of the result data by comparing the grouping set query's presentation**

Reference:

<https://cwiki.apache.org/confluence/display/Hive/Enhanced+Aggregation%2C+Cube%2C+Grouping+and+Rollup>

```
SELECT driverId, eventType, COUNT(*) AS occurrence
FROM csc534.tdend2_truck_event
GROUP BY driverId, eventType
WITH ROLLUP;
```

```
tdend2@node00:~
hive> SELECT driverId, eventType, COUNT(*) AS occurrence
> FROM csc534.tdend2_truck_event
> GROUP BY driverId, eventType
> WITH ROLLUP;
Query ID = tdend2_20241009181236_d45776da-5356-4265-a75c-a9e1a94ab82a
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
```

#### • The first 10 lines of query results only:

```
tdend2@node00:~
Total MapReduce CPU Time Spent: 8 seconds 170 msec
OK
NULL      NULL      17075
10         NULL      730
10         Normal   727
10         Overspeed    2
10         Unsafe tail distance  1
11         NULL      686
11         Lane Departure  4
11         Normal   676
11         Overspeed    1
11         Unsafe following distance  1
11         Unsafe tail distance  4
```

From the output, it can be observed that each row represents the 'driverid' and its corresponding 'event type' along with the sum of event\_count for that combination. The difference in the results of grouping set and ROLL UP is that the first row in the result of ROLL UP is not obtained with GROUPING SET. The difference in the number of resultant rows of GROUPING SETS and ROLL UP is less.

*Hence, the difference between GROUPING SETS and ROLLUP is clearly visible from the first half of the result especially the first 10 rows and found in the number of resultant rows and the order of it is GROUPING SETS < ROLL UP.*

## Cube

Write the grouping set query using CUBE. Run and show the results.

Explain the presentation of the result data by comparing both the grouping set query and the ROLLUP query's presentation.

- Replaced GROUPING SETS (...) with 'WITH CUBE

Reference:

<https://cwiki.apache.org/confluence/display/Hive/Enhanced+Aggregation%2C+Cube%2C+Grouping+and+Rollup>

```
SELECT driverId, eventType, count(*) AS occurrence
FROM csc534.tdend2_truck_event
GROUP BY driverId, eventType
WITH CUBE;
```

```
tdend2@node00:~
hive> SELECT driverId, eventType, count(*) AS occurrence
> FROM csc534.tdend2_truck_event
> GROUP BY driverId, eventType
> WITH CUBE;
Query ID = tdend2_20241009181704_530db350-d732-46de-ab16-cabcc08f2490
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
```

- The first 10 lines of query results only.

```
tdend2@node00:~
Total MapReduce CPU Time Spent: 7 seconds 790 msec
OK
NULL      NULL      17075
NULL      Lane Departure  11
NULL      Normal    17041
NULL      Overspeed     9
NULL      Unsafe following distance  7
NULL      Unsafe tail distance  7
10        NULL      730
10        Normal    727
10        Overspeed     2
10        Unsafe tail distance  1
```

**GROUPING SETS:** Manually defines the exact sets of groupings you want.

**ROLLUP:** A hierarchical aggregation, showing subtotals and grand totals.

**CUBE:** Includes all possible combinations of the groupings, leading to more rows than ROLLUP or GROUPING SETS.

From the above cube output, each row represents 'driverid' and corresponding 'event type' along with sum of event-count for that combination. NULL values in the output represents totals or subtotals for combination where the corresponding column value is not included in the grouping set. The first row in the result of ROLL UP is not obtained with GROUPING SET but it is the common row for CUBE and ROLL UP in addition to other rows which are available in all these

results. In addition to this first row, other 5 rows with driver id 'NULL' are extra rows obtained in the result of CUBE.

**The difference in the number of resultant rows of GROUPING SETS and ROLL UP is not much but CUBE operation resulted in more number of rows than these operations as it includes all possible combinations of the groupings, leading to more rows than ROLLUP or GROUPING SETS.**

**Hence, the difference among GROUPING SETS, ROLLUP and CUBE is clearly visible from the first half of the result especially the first 10 rows and found in the number of resultant rows and the order of it is GROUPING SETS < ROLL UP < CUBE.**

Here we use a machine learning competition dataset from Kaggle, "Titanic: Machine Learning from Disaster". See the detail <https://www.kaggle.com/c/titanic>.

We will use simple descriptive statistics to understand and analyze the data using HiveQL.  
**Found the dataset in the cluster:**

```
tdend2@node00:~  
[tdend2@node00 ~]$ head /home/data/CSC534BDA/datasets/Titanic/titanic.csv  
PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked  
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S  
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C  
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S  
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S  
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S  
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q  
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S  
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S  
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S  
[tdend2@node00 ~]$
```

## Step1

**Load the data to your user space in HDFS**

**Created directory as follows to copy data from local file system to hdfs:**

**hdfs dfs -mkdir /user/tdend2/titanic**

**hdfs dfs -copyFromLocal /home/data/CSC534BDA/datasets/Titanic/titanic.csv  
/user/tdend2/titanic**

```
tdend2@node00:~  
[tdend2@node00 ~]$ hdfs dfs -mkdir /user/tdend2/titanic  
  
tdend2@node00:~  
[tdend2@node00 ~]$ hdfs dfs -ls /user/tdend2/titanic/  
Found 1 items  
-rw-r--r--  3 tdend2 hadoop      61194 2024-10-11 11:34 /user/tdend2/titanic/titanic.csv  
[tdend2@node00 ~]$
```

Also, file can be uploaded using the below command to the desired directory:

**hdfs dfs -put /home/data/CSC534BDA/datasets/Titanic/titanic.csv /user/tdend2/titanic/**

Directory is created successfully and loaded contents of titanic as shown below:

```
tdend2@node00:~  
[tdend2@node00 ~]$ hdfs dfs -ls /user/tdend2/titanic/  
Found 1 items  
-rw-r--r-- 3 tdend2 hadoop 61194 2024-10-10 15:17 /user/tdend2/titanic/titanic.csv  
[tdend2@node00 ~]$ hdfs dfs -cat /user/tdend2/titanic/titanic.csv | head  
PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked  
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S  
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C  
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S  
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S  
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S  
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q  
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S  
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S  
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S  
[tdend2@node00 ~]$
```

As seen above, 12 columns exist in titanic.csv, which are as follows:

PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked

## Step2

### Create a Hive table

- We may need to know how to handle double quotes.

PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked

In hive shell create hive table 'titanic' using the above column names as follows:

```
CREATE TABLE csc534.tdend2_titanic (  
    PassengerId INT,  
    Survived INT,  
    Pclass INT,  
    Name STRING,  
    Sex STRING,  
    Age FLOAT,  
    SibSp INT,  
    Parch INT,  
    Ticket STRING,  
    Fare FLOAT,  
    Cabin STRING,  
    Embarked STRING  
)  
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'  
WITH SERDEPROPERTIES (  
    "separatorChar" = ",",  
    "quoteChar" = "\"" )  
STORED AS TEXTFILE;
```

```
tdend2@node00:~  
hive> CREATE TABLE csc534.tdend2_titanic (  
  > PassengerId INT,  
  > Survived INT,  
  > Pclass INT,  
  > Name STRING,  
  > Sex STRING,  
  > Age FLOAT,  
  > SibSp INT,  
  > Parch INT,  
  > Ticket STRING,  
  > Fare FLOAT,  
  > Cabin STRING,  
  > Embarked STRING  
  > )  
  > ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'  
  > WITH SERDEPROPERTIES (  
  > "separatorChar" = ",",  
  > "quoteChar" = "\""  
  > )  
  > STORED AS TEXTFILE;  
OK  
Time taken: 0.691 seconds  
hive>
```

### Step3

#### Load the data to Hive

We assumed the dataset is big, so loaded the data from HDFS, not the local filesystem.

```
LOAD DATA INPATH '/user/tdend2/titanic/titanic.csv' INTO TABLE csc534.tdend2_titanic;
```

```
tdend2@node00:~  
hive> LOAD DATA INPATH '/user/tdend2/titanic/titanic.csv' INTO TABLE csc534.tdend2_titanic;  
Loading data to table csc534.tdend2_titanic  
OK  
Time taken: 1.048 seconds  
hive>
```

**Verification:** if loading was successful, comparing the line/row count of both source and target.

```
select count(*) from csc534.tdend2_titanic;
```

```

tdend2@node00:~
hive> select count(*) from csc534.tdend2_titanic;
Query ID = tdend2_20241011114635_46d90142-3baf-4891-947c-8db9ca4c2ef9
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
24/10/11 11:46:35 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
24/10/11 11:46:35 INFO client.RMProxy: Connecting to ResourceManager at node00.sun/10.0.0.10:8032
Starting Job = job_1722897143033_1446, Tracking URL = http://node00.sun:8088/proxy/application_17
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/lib/hadoop/bin/hadoop job -
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-10-11 11:46:45,330 Stage-1 map = 0%, reduce = 0%
2024-10-11 11:46:50,475 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.43 sec
2024-10-11 11:46:58,672 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.43 sec
MapReduce Total cumulative CPU time: 5 seconds 890 msec
Ended Job = job_1722897143033_1446
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.89 sec HDFS Read: 88918 HDFS Write: 103 HD
Total MapReduce CPU Time Spent: 5 seconds 890 msec
OK
891
Time taken: 25.121 seconds, Fetched: 1 row(s)
hive>

```

**Row count is 891**

*This matched with line count of titanic.csv(included header) which is available in local linux file system as shown below:*

**Row count of titanic.csv available in localfile system:**

`wc -l /home/data/CSC534BDA/datasets/Titanic/titanic.csv`

```

tdend2@node00:~
[tdend2@node00 ~]$ wc -l /home/data/CSC534BDA/datasets/Titanic/titanic.csv
892 /home/data/CSC534BDA/datasets/Titanic/titanic.csv
[tdend2@node00 ~]$

```

**892 lines are found, excluding header row of csv file, row count is 891.**

- If we drop(delete) your table, we should reload the data to HDFS because the LOAD command moves the data in space of HDFS to Hive storage of HDFS.



Write analytical queries (descriptive statistics) to answer the below questions.

#### Step4

Calculate the survivor counts and survival/death rates.

Survivor count and survival/death rates?

```
SELECT
    survived,
    COUNT(*) AS total,
    COUNT(*) / CAST(SUM(COUNT(*)) OVER () AS DOUBLE) AS survival_death_rate
FROM csc534.tdend2_titanic
GROUP BY survived;
```

```
tdend2@node00:~
hive> SELECT
>     survived,
>     COUNT(*) AS total,
>     COUNT(*) / CAST(SUM(COUNT(*)) OVER () AS DOUBLE) AS survival_death_rate
> FROM csc534.tdend2_titanic
> GROUP BY survived;
Query ID = tdend2_20241012114058_73d1d7cb-d870-4602-96b8-85cbaf4f4faf
```

Result:

```
Total MapReduce CPU Time Spent: 12 seconds 880 msec
OK
survived      total      survival_death_rate
1             342      0.3838383838383838
0             549      0.6161616161616161
Time taken: 48.792 seconds, Fetched: 2 row(s)
hive>
```

**0 -> non-survivors from survived column: death count is 549 and accounted to 61.62%**

**1 -> survivors from survived column: survivors count is 342 and accounted to 38.38%**

From the above, it can be concluded that death rate of titanic passengers is more than those who survived(survival rate).

#### Step5 (10pts)

Calculate the survivor counts and survival/death rates by survived

and Sex. Note: The sum of survival rates should be 1 or 100%.

```
SELECT
    survived,
    sex,
    COUNT(*) AS total,
    COUNT(*) / CAST(SUM(COUNT(*)) OVER (PARTITION BY Survived) AS DOUBLE) AS
survival_death_rate
FROM csc534.tdend2_titanic
GROUP BY survived, sex;
```

```
hive> SELECT
>     survived,
>     sex,
>     COUNT(*) AS total,
>     COUNT(*) / CAST(SUM(COUNT(*)) OVER (PARTITION BY Survived) AS DOUBLE) AS survival_death_rate
> FROM csc534.tdend2_titanic
> GROUP BY survived, sex;
Query ID = tdend2_20241012114511_1f01965b-140b-4e03-8e7d-b1ed0a418fdd
```

The output is as below:

```
Total MapReduce CPU Time Spent: 7 seconds 990 msec
OK
survived      sex      total      survival_death_rate
0      female  81      0.14754098360655737
0      male    468     0.8524590163934426
1      female  233     0.6812865497076024
1      male    109     0.31871345029239767
Time taken: 24.058 seconds, Fetched: 4 row(s)
hive>
```

0-> non-survivor(deceased)

1-> survivor

**Survivor counts and survival/death rates by Sex:**

0-> non-survivor(deceased) count is 81 and rate among female is 14.75%

0-> non-survivor(deceased) count is 468 and rate among male is 85.25%

1-> survivor count is 233 and rate among female is 68.13%

1-> survivor count is 109 and rate among male is 31.87%

From the above, it can be concluded that **female sex(68.13%) is higher in survival rate (not death rate) when compared to male survivors (31.87%)**

## Step6

**Write a query to produce similar results of Step 4 and Step 5 Combined.**

```
SELECT
    survived,
    sex,
    COUNT(*) AS total,
    COUNT(*) / CAST(SUM(COUNT(*)) OVER (PARTITION BY Survived) AS DOUBLE) AS
survival_death_rate
FROM csc534.tdend2_titanic
GROUP BY survived, sex WITH CUBE;
```

```
hive> SELECT
>     survived,
>     sex,
>     COUNT(*) AS total,
>     COUNT(*) / CAST(SUM(COUNT(*)) OVER (PARTITION BY Survived) AS DOUBLE) AS survival_death_rate
> FROM csc534.tdend2_titanic
> GROUP BY survived, sex WITH CUBE;
Query ID = tdend2_20241012115344_28c94442-242b-4426-8caf-618a2d3c7823
```

The output is as below:

```
Total MapReduce CPU Time Spent: 7 seconds 980 msec
OK
survived      sex      total      survival_death_rate
NULL          NULL      891        0.5
NULL          female    314        0.17620650953984288
NULL          male      577        0.32379349046015715
0             NULL      549        0.5
0             female    81         0.07377049180327869
0             male      468        0.4262295081967213
1             NULL      342        0.5
1             female    233        0.3406432748538012
1             male      109        0.15935672514619884
Time taken: 25.062 seconds, Fetched: 9 row(s)
hive> █
```

=====THE END=====